

On Datacenter Performance and Climate-Impact with Multi-Model Simulation and Analysis

Abstract—Datacenters are vital for the digital society but represent a considerable fraction of global energy consumption. To improve their sustainability and performance when demand is foreseen to increase, we envision simulators will become primary decision-making tools. However, unlike other fields focusing on key societal infrastructure such as waterworks and mass transit, datacenter simulators cannot yet combine multiple, independent models into their operation. Addressing this challenge, in this work we propose M3SA, a datacenter simulation and analysis framework that uses discrete-event simulation to predict, per model and then combined into a meta-model, the impact on climate and performance of various realistic datacenter conditions. We design an architecture for simulating multiple concurrent models, a technique to integrate the results of multiple models into a meta-model, and a procedure to evaluate the accuracy of the meta-model. Through experiments with a prototype, we show that (i) M3SA can be used to reproduce peer-reviewed experiments, and enhance their output with more diverse metrics and more detailed analysis; (ii) M3SA can be configured with a variety of realistic parameters, such as diverse workload traces (using Grid Workloads Archive data in our experiments), and energy production data such as carbon intensity over time and location (gCO₂/kWh data across all EU regions, from ENTSO-E); (iii) M3SA enables various types of what-if and how-to analysis, such as how to configure CO₂-aware migration over yearly energy-production patterns. M3SA has been integrated into the open-source software simulator OpenDC and is available on GitHub¹.

Index Terms—datacenters, multi-model simulation, performance, energy utilization, climate impact, sustainability, OpenDC

I. INTRODUCTION

Our society and economy are increasingly dependent on digital services. Correspondingly, new datacenters are being built, and existing datacenters are scaled up [1]–[4]. Before building or expanding, operators use simulators to project infrastructure performance and climate impact [5]–[7]. For example, in the EU research project Graph Massiziver, simulators predict speedup, failure cost, energy consumption, and CO₂ emissions for massive-scale infrastructure [8]. Although datacenter simulation remains challenging in general, for example due to the unpredictable behavior of datacenter software and hardware ecosystems [9], [10] and the continuous emergence of new devices and applications [3], [4], in this work we focus on the model at the core of simulation. Current simulators, such as the open-source OpenDC [11], CloudSim [12], [13], and SimGrid [14], [15], already support *single-model simulation*—the simulation is based on a single integral model, where each component or phenomenon is modeled with a partial

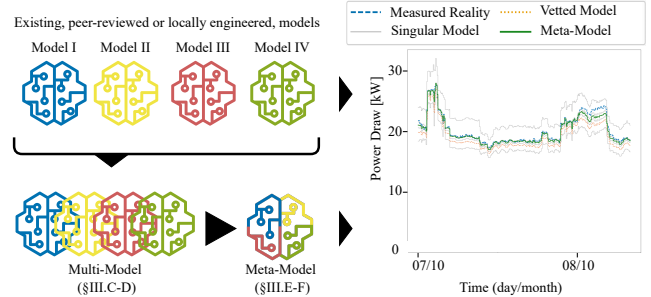


Fig. 1. M3SA Multi-Model simulation uses multiple models simultaneously. M3SA Meta-Model integrates all results. Figure 9 depicts a specific scenario.

model, e.g., one for CPUs and one for the storage subsystem. The partial models are conceptually independent, which facilitates scaling the simulation. When each partial model fails to address current needs, new models are developed, tuned, or trained for specific scenarios, at significant cost and time [16]–[18]. In turn, single integral models are replaced by rewiring existing partial models. However, as this work showcases by quantifying single-model error in Section IV-B, simulations based on *single* models are prone to errors when encountering edge cases and new scenarios. Addressing this gap, in this work we design, prototype, and evaluate the M3SA (read as [mesa]) framework for *multi-model* (and *meta-model*) simulation and analysis of datacenters.

Figure 1 illustrates our approach, as a high-level abstraction (left) and through a real-world experiment (right, sample from §IV-B). Here, M3SA considers four existing single integral models, Models I-IV, which are peer-reviewed or engineered by the datacenter operator. Then, M3SA conducts multi-model simulation, automating the independent execution of each model. Last, M3SA integrates the results into the Meta-Model. Figure 1 (right) showcases the ability of the Meta-Model to alleviate singular model biases and better align with the measured reality; the Meta-Model delivers even better accuracy than a state-of-the-art, peer-reviewed model [16].

Albeit essential, single integral (*singular*) models are insufficient for reliable predictions. A singular model makes accurate predictions in the context it has been trained for and often fails to handle real-life edge-case scenarios [11], [19], [20]. Moreover, singular models are calibrated to ideal scenarios and may lack adaptability to different *operational phenomena*, such as sudden spikes in demand, performance variability, equipment failures, and power outages. Such scenarios can lead to significant discrepancies between model-prediction

¹URL for double-blind review: <https://anonymous.4open.science/r/m3sa/>

and reality, leading to costly downtimes [21], impactful outages [22], and datacenter shutdowns [23]. Although model tuning or re-training can take place, the resulting models could become prone to errors with the occurrence of new operational phenomena or with changes in the baseline scenarios they represent.

In this work, we posit that combining multiple models and contrasting their predictions for the same simulated scenario, can retain the strengths of each model and alleviate the biases of each model. The community leverages diverse modeling approaches trained or validated on diverse datasets [11], [17], [18], [24]–[26]—each an opportunity to include multi-model simulation. However, none of the open-source datacenter simulators currently supports multi-model simulation, including the highly cited and much-used CloudSim [12], [13], SimGrid [14], [15], OpenDC [11], and GDCSim [27]. One of the possible drawbacks of multi-model simulation, which we test for explicitly in this work, is the computational overhead in simulation, over traditional -model approaches; we quantify this overhead in Section IV-B.

Simulating using multiple models is a novelty in computer systems; conducting multi-model simulations for datacenters is also innovative beyond computer systems science. Coarse-grained, large-scale models in weather and climate simulation [20], [28] and detailed, small-scale models in ecology [29] already use simulation instruments that rely on multiple prediction models, calibrated and adjusted for the needs of their scientific field. In contrast, to obtain meaningful predictions about complex datacenters ecosystems, medium-scale models must combine higher-level abstraction with detailed operational models for specific devices and applications, and address various operational phenomena. This is fundamentally different than other-scale sciences [30].

In this work, we address the open challenge of *leveraging multiple models and integrating their results into a Meta-Model* by designing, implementing, and evaluating a first-of-its-kind tool, M3SA, in complex, realistic, trace-driven simulations, with a three-fold contribution:

- C1 (Design)** We design M3SA, the first multi- and meta-model simulator for datacenters (Section III). We formulate functional and non-functional requirements. We propose an overall architecture for Multi- and Meta-simulation of datacenters. We further design the Multi- and Meta-Model components to work with generic discrete-event simulators, and propose assisted configuration and evaluation (validation) processes. Overall, M3SA allows the same component to be modelled in different ways and systematically combines predictions of distinct, independently run, singular models.
- C2 (Experimentation)** We prototype M3SA and conduct three large-scale, trace-driven datacenter experiments (Section IV). We embed M3SA into OpenDC [11], a peer-reviewed, state-of-the-art datacenter simulator, and, to showcase its practical capabilities, equip it with 8 commonly used energy models for the computational component. We conduct with M3SA three experiments,

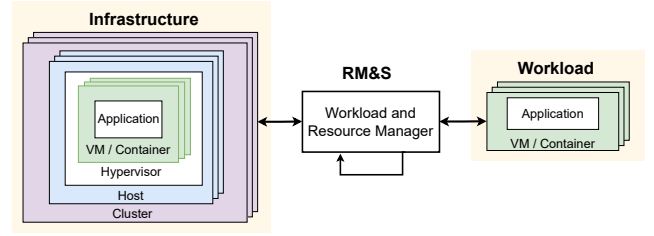


Fig. 2. Simulation enables fine-grained infrastructure operational monitoring used by stakeholders in decision-making processes. Adapted from [35].

to: (i) Reproduce a peer-reviewed, open-science datacenter experiment [16] (Section IV-B); (ii) Analyze and compare the impact of workload kind—scientific [31] vs. business-critical [32] on datacenter performance and CO₂-emissions (Section IV-C); and (iii) Analyze the impact of EU-region on the performance and CO₂-emissions of migrated workload (Section IV-D). Overall, M3SA enables the analysis of both datacenter performance and CO₂-emissions related to resource use, across diverse workloads and geographical locations, with improved accuracy and explainability over -model simulations. Our Multi-Model approach can enhance prediction accuracy and explainability over commonly used models and on-par with a state-of-the-art, hand-tuned model (§IV-B) and supports scenarios for which no open-source hand-tuned model exists (§IV-C and §IV-D). The M3SA tool, assembling the Multi- and Meta-Model, can simulate years of real-world operation within minutes (e.g., ≈ 4 years in ≈ 11 minutes), on a regular user machine (e.g., Macbook M2 Pro) (§IV-B).

- C3 (Open science)** We contribute to open science all the software and data [33] created in this project. We engineer the M3SA software and integrate it into OpenDC, which is based on the popular Kotlin and Python languages [34]. For the last experiment, we collect public CO₂-emissions data across all EU countries from ENTSO-E, align the timestamps, and compute carbon-intensity data from it, and release the resulting FAIR dataset. Last, we make available the complete reproducibility capsule: <https://anonymous.4open.science/r/m3sa/>.

II. BACKGROUND: SIMULATION AND SYSTEM MODEL

We present background on datacenter simulators and a basic system model, illustrated by Figure 2. Combining the two, we present elements of prediction models and metrics widely used in the community and useful for understanding this work.

“Simulation is defined as the imitation of the operation of a system or real-world process over time, and in many cases, manufacturing provides one of the most important applications of simulation” [36]. Simulation provides datacenter stakeholders with operational insights into the behavior of infrastructure running under a given workload.

Figure 2 illustrates our basic system model, which is an abstraction of many real-world operations. *Workloads* contain tasks operated on *infrastructure*, which includes physical

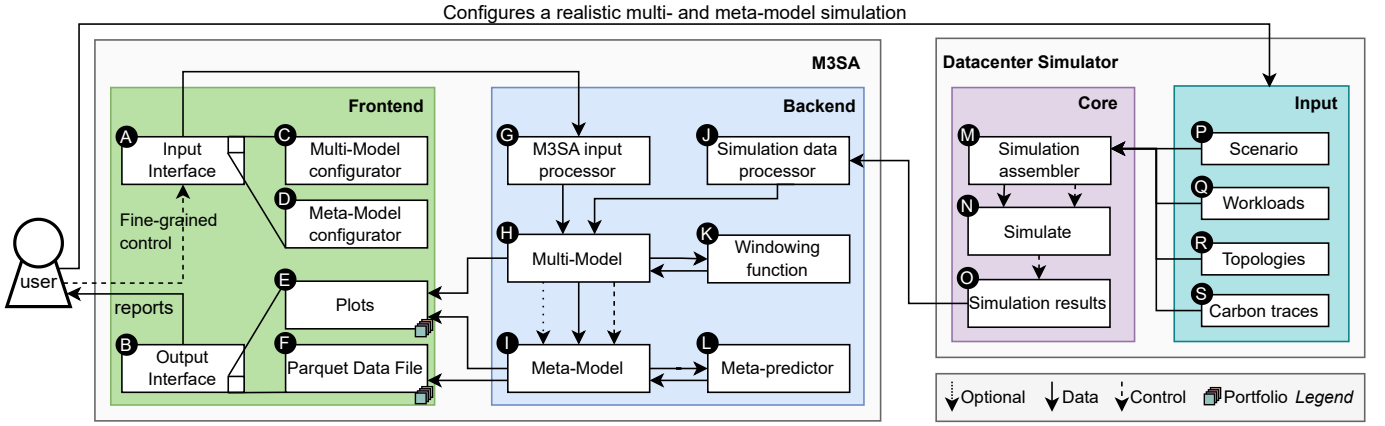


Fig. 3. Architecture overview of M3SA, a Multi- and Meta- Model Simulation Analyzer, integrated with a black-boxed simulator.

machines, virtual machines (VM), and containers; a complex *resource management and scheduling (RM&S)* system is continuously matching workloads to infrastructure, subject to optimization goals such as energy use (power draw) and CO₂-emission reduction. *Traces* represent recordings of real-world events, capturing detailed operational data of the infrastructure under given workload(s); traces provide a granular view of resource usage, essential for driving simulations or replaying real-world scenarios.

Predictive models in large-scale computer systems are empirical prediction systems that analyze, combine, and compute various input elements (often, traces) to produce fine-grained output predictions [37, §1]. We use such models to predict how real-world workloads run on ICT infrastructure, under various operational phenomena, as modeled by users. In this work, we use three types of models: (i) *Energy models* predict the amount of energy the datacenter consumes from the power grid. (ii) *CO₂ emission models* predict the carbon intensity (CO₂ emitted per unit of energy, measured in gCO_2) of the datacenter under the simulated workload. (iii) *Failure models* simulate task failures according to a failure trace. Without such traces, failure models generate task failures with arbitrary but realistic occurrence timestamps and recovery periods. Our technical report attached to this work [38] details each model and its analytical model.

We use *metrics* for system monitoring and accuracy measurement. For system monitoring, we use two main metrics: *energy consumption* and *carbon intensity*, measured in international system units. We define and measure the models' accuracy against real-world data (i.e., measured reality) using *Mean Percentage Absolute Error (MAPE)*, expanded in Section III-F.

III. DESIGN OF M3SA: A TOOL FOR MULTI- AND META-MODEL SIMULATION AND ANALYSIS

In this section, we synthesize requirements and design a novel approach for datacenter simulation. We propose M3SA, a system that enables simulation using *multiple* models.

A. Requirement Analysis

(FR1) Leverage multiple models into a unified tool: The system must enable combining multiple (existing) models into

a unified tool. It should enable visualizing their predictions, independently or combined, thus increasing the explainability of predictions, and pointing out idiosyncratic errors or biases. **(FR2) Predict by integrating multiple models:** A *Meta-Model*, able to aggregate singular-models' predictions, at a user-selected granularity, to alleviate errors and biases of idiosyncratic models, thus strengthening the credibility of the predictions.

(FR3) Open-source tool: The system should be open-sourced. The system should contribute to open science, both as coupled with a peer-reviewed, vetted, and state-of-the-art simulator and decoupled, as an independent scientific tool.

(NFR1) Provide in-meeting or same-day simulation results: Multiple models can be run in parallel, reducing the simulation time to a matter of minutes or less². The analysis part of the system should only add a fraction to the simulation time itself, and not more than double it.

(NFR2) Highly accurate predictions: The Meta-Model should predict with higher accuracy than the average singular model employed in the simulation. Equivalently, the error rate of the Meta-Model should be lower than the error rate of the average model.

(NFR3) Operation at real-world scale: All FRs and NFRs should be met for *at least* 8 singular models, on input traces of *at least* 200,000 samples per model, equivalent to 2 years at the industry/science de facto sampling rate of 5 minutes [11], [35]. Although a more general aim would be to support as many models as possible, as the experimental results in §IV-B indicate, already 8 such models, combined, can deliver good results when compared with state-of-the-art hand-tuned models.

B. Overview of the M3SA Architecture

We design M3SA to be capable of operating coupled or decoupled from a datacenter simulator. Figure 3 depicts an overview of the system's architecture, in which M3SA extends a black-boxed simulator **(FR3)**. Throughout this paper, we couple M3SA with OpenDC, a peer-reviewed, open-source,

²Complex datacenter experiments using OpenDC can be used as a reference point. They are reported to take seconds to minutes [11].

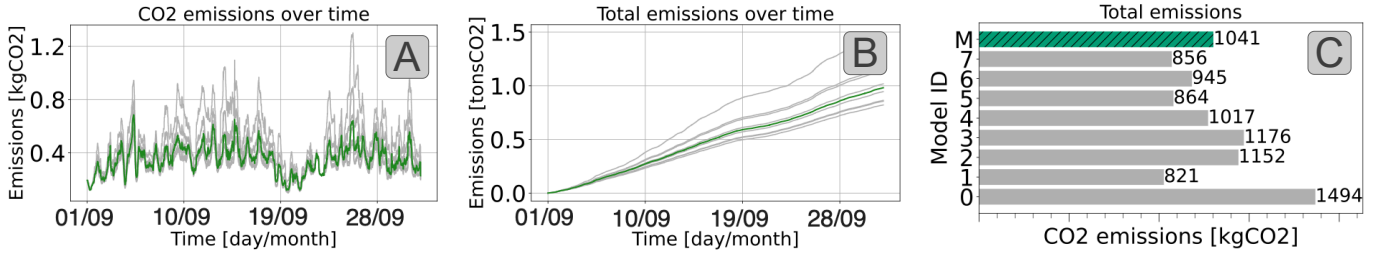


Fig. 4. Sample M3SA plots. We represent singular models in gray and Meta-Models in green. In sub-plots A and B: vertical axis depicts the simulated metric, horizontal axis depicts time. In sub-plot C: vertical axis depicts model identifier, horizontal depicts cumulative totals.

discrete-event simulator with simple interfaces, and over 5 years of development and operation [1], [11].

M3SA adopts the architecture we define as “Simulate First, Compute Later” (SFCL), where M3SA operates only after the simulator successfully predicts and the *simulation results* (O) are available. This architecture couples M3SA to the simulator as an external layer. We also considered the alternative architecture of “Compute While Simulating” (CWS), where M3SA would be deeply embedded into the simulation infrastructure. Comparing SFCL and CWS architectures, we observe a tradeoff between the depth of M3SA embedding (FR3) and the performance of the overall system (NFR1) - the more embedded M3SA is, the better the overall performance, yet the more complex it becomes to port M3SA from one simulator to another. We choose the SFCL architecture, thus allowing simple adoption of M3SA; to meet the performance requirements, we focus on reducing redundancies and following good engineering practices.

In the SFCL architecture, the user interacts with the system through the *Input Interface* (component A) and *Output Interface* (B) interfaces (FR3). The M3SA process begins from the input, where the users configure the *Multi-Model* (C) and the *Meta-Model* (D). The simulation process, triggered and controlled by the M3SA backend, happens between M-S: the system sets up a simulation based on user-input, simulates, and centralizes predictions. The simulation block M-S corresponds to how discrete-event simulators commonly used in the field actually operate, and closely matches the OpenDC [11] and CloudSim [13] architectures; in particular, the simulation assembler (M) is where singular models are typically defined in current experiments.

Further, the M3SA backend retrieves the predictions made by the simulation infrastructure, one per singular model, and assembles the *Multi-Model* (FR1). To avoid redundant computation, the *Meta-Model* predicts using already leveraged and processed data from the *Multi-Model* (FR2), (NFR1). Lastly, the system plots the predictions of the *Multi-Model* and *Meta-Model* (E) and outputs the *Meta-Model* predictions to a lossless-compressed data file (F).

C. Design of the Multi-Model Component

In this section, we define explainability, and elaborate on the *Multi-Model* and *Plots* components (Figure 3 H and E, respectively). The M3SA multi-model approach helps compare the outputs singular models, thus improving users’ understand-

ing of the predictions and increasing explainability (FR1). It also enables meta-modeling, as detailed in Section III-E.

We define *explainability* as the user’s understanding of the possible behavior, limitations, and biases of the system under test, as predicted by relevant, peer-reviewed or locally vetted, models. Better explainability is achieved when multiple models are available, and when at least the ranges of acceptable predictions, and also the detailed singular predictions, are available and contrasted across the models, enabling specialists to better reason about the problem³. Ideal explainability both matches measurable reality and links real-world values to predictions made by one or several models, thereby enabling specialists to explain wherever evidence exists about which model gives better predictions and how close those predictions are to instrumental observations (traces)—proxy to reality.

The *Multi-Model* component (H) leverages multiple models into a unified tool. Within this component, predictions of singular models, which M3SA has already produced through its Datacenter Simulator backend, are centralized; the user can select the form of data representation. To increase performance and alleviate redundancy, the system reads from the input data file only the metrics selected by the user (typically, at configuration time) and discards irrelevant data. The *Multi-Model* process then filters prediction data from singular models with a user-selected windowing function, which minimizes the dataset for better performance (NFR1) and also removes noise for better visual comprehension; we expand on the concept of windowing in Section III-D. After completion, the *Multi-Model* is ready for further operations, such as *Meta-Model* computation or plotting (FR2).

The *Plots* component (E) uses the *Multi-Model* results to produce, guided by typical user-given configuration, visual outcomes as in Figure 4. M3SA supports analyzing how metrics evolve over time, enabling the reader to analyze predictions at specific timestamps (Figure 4A), how quantities accumulate over time (Figure 4B), and aggregates such as sums (totals) over a part or the entire experiment (Figure 4C).

D. Analysis and Configuration of Multi-Model Window Size

To address the tradeoff between the detail of the visual comprehension (FR1) and the performance (NFR1) aspects of

³In many scientific domains, e.g., in weather and climate modeling [20], [28], multiple instrumental observations (traces) and multiple models are used to provide evidence, toward building sufficient confidence in specific outcomes.

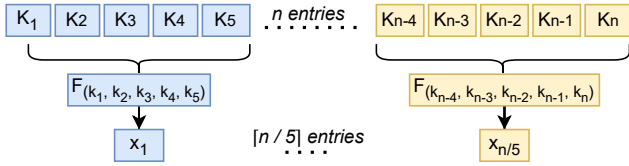


Fig. 5. Window size 5 applied on n entries, using aggregation function F .

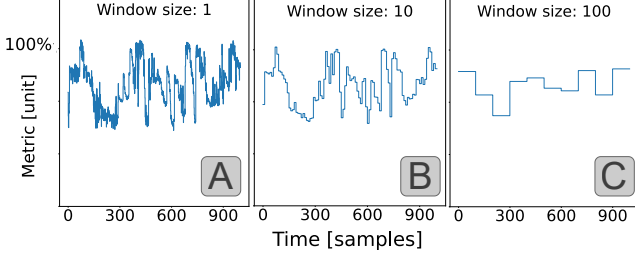


Fig. 6. Comparison of visual quality when windowing representative input data. Window sizes differ per sub-plot, in order, 1, 20, and 100. The horizontal axis represents time, in abstract units. The vertical axis represents the system response metric, from zero to maximum.

the multi-model, we equip the Multi-Model component with a windowing mechanism.

We define a *window* as an aggregation of chunks of data into a singular data entry, similar to a one-dimensional convolutional layer. Figure 5 illustrates how a window of size m (here, $m = 5$) is applied to n data entries. Each group of m elements is fed into an aggregation function F , leading to a final compressed output array of size $\lceil n/m \rceil$.

We now explore the trade-off between visual quality and performance. We conduct experiments with actual outputs produced by the simulator, for input between about 2,000 and 200,000 samples long; with various window sizes between 1, corresponding to retaining all data, and 1000, a large value; and without M3SA-level optimizations such as eliminating redundancies. (When optimizations are enabled, as indicated by the results in Section IV-B, the performance of M3SA becomes compatible with (NFR1).) Figure 6 presents a representative selection of the results. As expected, the larger the window size, the sparser the filtered output and thus the coarser the visual detail. At a window size of 100, and beyond, the actual shape of the input is irretrievably lost. On the other side of the trade-off, the parsing and graphing sequence of our largest test-input ($n = 201,600$), when $m = 1$ requires, on a common, off-the-shelf, Apple M2 MacBook Pro, about 4.5 minutes (264.7 s). As the window size filters out data to persist and graph, the runtime of computation and graphing decreases as the window size increases, to just over 3 minutes at $m = 10$. From $m = 20$, additional effects, e.g., the working set becoming increasingly large, reverse this trend, and the time needed when $m = 1,000$ for the largest input already approaches the time taken without windowing (238.3 s). Based on these experimental results, we configure the Multi-Model with windowing-sizes of $m = 1$ (§IV-B) and of $m = 10$ (§IV-C, §IV-D) for this work.

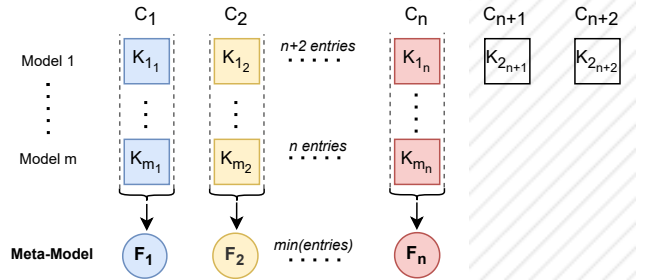


Fig. 7. The Meta-Model predictor, aggregating predictions of m models vertically, for each time-step C_1 to C_n . Steps C_{n+1} and C_{n+2} are discarded.

E. Design of Meta-Model component

In this section, we detail the *Meta-Model* component (Figure 3, ①). Overall, the Meta-Model uses an automated process to combine predictions made with multiple models (FR2) into a highly accurate result (NFR2).

The Meta-Model relies on the predictions made by multiple models. For performance reasons (NFR1), this process leverages processing results from the Multi-Model (see Section III-C), thus avoiding redundant computation. Upon initialization, the core of the Meta-Model uses the Meta-predictor (①), which combines other models' (windowed) predictions (FR2). After ① completes, the Meta-Model contains a one-dimensional set of predictions, prepared for output.

The output involves two independent processes, for usability (FR1), performance (NFR1), and portability (FR3) reasons. First, the system compresses the predictions losslessly, into a Parquet data file, which the user can process outside of the M3SA system (⑤); we use Parquet due to the scalability, efficient storage, and cross-system compatibility (and thus portability) of this storage format [39]. Second, the Meta-Model is plotted (⑤) following user inputs (A).

We define *robustness* as the confidence that predictions perfectly and consistently match the measured reality. Although the *accuracy* of the Meta-Model is directly dependent on the accuracy of singular models (①) and on the aggregation function (K), the robustness of the Meta-Model accuracy should intuitively be higher than of any singular model (NFR2), because the Meta-Model can mitigate the errors of singular models when it combines the results across multiple models. This can be expressed and empirically measured for specific inputs, e.g., the MAPE metric introduced in Section III-F.

To achieve high accuracy, the Meta-Model component uses the Meta-predictor (①). The Meta-predictor receives a set of predictions over time, one per input model, with time divided in equal chunks (steps). As Figure 7 illustrates, the Meta-predictor first aligns the steps across all models, for example, removing entries when too few models provide predictions, and aggregates the remainder. In the figure, the rows represent the predictions of each of the m models, and the columns represent the aligned time-steps. Model 1 provides additional predictions, for time steps C_{n+1} and C_{n+2} , which are discarded by the Meta-predictor because not enough other models also provide predictions for these steps.

Following alignment, the Meta-predictor applies an aggregation function. In the figure, the values stored in each column C_k are aggregated for each time-step, using a (user-configurable) function F_k . In Figure 7, the aggregation function is applied vertically for each column, which is convenient also for performance. However, the concept also allows for diverse spatial combinations and further, more complex but model-order-free, approaches. We leave for future work the exploration of such aggregators, using mathematical, statistical, and machine-learning techniques.

For the aggregating function F , we explored for this work two computationally light, commonly used aggregation functions, (arithmetic) *mean*, which is used in ecology [29] and meteorology [20] but has high sensitivity to outliers, and *median*, which alleviates the outlier sensitivity of the mean but is more computationally expensive.

The proposed Meta-Model design poses two challenges. Firstly, timestamp misalignments may occur due to non-deterministic simulation infrastructure. We address this challenge by selecting a datacenter simulator that establishes the same initial sampling timestamp for all models and the same sampling rate of the infrastructure. Secondly, singular models may output prediction datasets of different sizes (i.e., longer or shorter virtual infrastructure execution times), as we illustrate in Figure 7, for models 1 and m; this difference may occur due to various reasons, such as scheduling algorithms that may prioritize tasks differently, or the simulation of potential infrastructure failures. To address this challenge, the meta-predictor uses only the minimum number of entries per model in the computation process.

This design of the Meta-Model already allows for reducing the error rates of peer-reviewed, state-of-the-art models (§IV-B). This design, yet still powerful could be further explored by employing machine learning techniques in aggregating models, or by dynamically, on-the-fly, allocating weights to the singular models, or employing Multi-Criteria Decision Analysis (MCDA), which are widely used decision-making processes in other-scale sciences [40]–[42]. Although this is beyond the purpose of this research, we envision future work which could analyze such techniques, trade-offs performance-accuracy, and advantages of more complex, yet more computationally heavy, aggregation techniques.

F. Evaluation of Meta-Model Accuracy

To quantify robustness (see Section III-E) in M3SA (**NFR2**), we use the Mean Absolute Percentage Error (MAPE), a widely used relative-error metric [16], [43], [44]. MAPE equally penalizes positive and negative errors and is calculated using Equation (1), where n is the number of samples, R is real-world data, S is simulation data, i is the sample index:

$$MAPE [\%] = \frac{1}{n} \sum_{i=0}^n \left| \frac{R_i - S_i}{R_i} \right| \times 100 \quad (1)$$

Anticipating the need for further accuracy evaluation (**NFR2**), we design M3SA to be extensible to other

metrics such as Normalized Absolute Difference (NAD) [16], Root of the Mean Square Error (RMSE) [45], and Mean Average Error (MAE) [46], [47].

IV. COMPLEX, TRACE-BASED EXPERIMENTS WITH M3SA

In this section, we conduct diverse experiments, evaluating M3SA and simultaneously providing new insights into datacenter operation. To setup our experiments (Section IV-A), we first implement a prototype of M3SA, interface it with OpenDC, and provide the system open-source (**NFR2**). In Section IV-B, we reproduce a peer-reviewed experiment with M3SA capabilities and analyze M3SA against established requirements. In Section IV-C, we run an experiment using two fundamentally different traces, analyze the results, and evaluate the versatility of M3SA. Lastly, we evaluate M3SA on CO2-aware migration over energy-production patterns. We focus here on the main findings and further analyze the experiments in a technical report [38].

Overall, the experiments support four main findings (MFs):

(**MF1**) Overall, the Multi-Model can increase the explainability of the simulation results, and help in identifying biases and errors of singular models; the Meta-Model can enhance prediction accuracy against singular models, robustly. Specifically, the Meta-Model can reduce the error rate of the average singular simulation model by about 50% (MAPE decreases from 7.59% to 3.81%) with under 20% computational overhead (Section IV-B).

(**MF2**) M3SA can simulate 1 (or 2) years of datacenter operation in 5 (7) minutes, of which under 0.5 (1) minutes are M3SA overhead (Section IV-B).

(**MF3**) Supporting various C-level decisions, M3SA can simulate business-critical and scientific workloads in scenarios including machine failures (Section IV-C).

(**MF4**) Adding to MF3, M3SA can help in CO2-aware workload migration and scheduling processes, with decisions that can help reduce CO2 emissions by 97.5% compared to the average location, and by 11% compared to the best but single location (Section IV-D).

A. Experiment Setup

We summarize our experimental design in Table I, and detail the infrastructure (systems), workloads, and carbon traces in Table II, Table III, and Table IV, respectively. We detail the most important aspects:

Workload traces (all public): SURF-22 is a scientific workload trace used in peer-reviewed experiments [16]. SURF-22 was traced in one of the largest HPC facilities at SURF, the Netherlands. It traces scientific jobs run in production, batch, with an average duration of 39.52 CPU-hours, and also detailed energy use over the same period.

Marconi-22 is one of the most detailed open-source workload traces currently available. It traces scientific jobs, and also multiple operational layers and the electrical network of Marconi HPC, a powerful supercomputing facility hosted by CINECA in Italy [31].

TABLE I

EXPERIMENT CONFIGURATIONS. (WT = WORKLOAD TRACE, CT = CO2 TRACE, PM = PERFORMANCE METRICS, SM = SUSTAINABILITY METRICS, IM = INDIVIDUAL MODELS, MW = MULTI-MODEL WINDOW SIZE, MF = META-MODEL AGGREGATION FUNCTION, SUO = SYSTEM UNDER OBSERVATION.)

Section	Focus	Input		Output					Failures	SUO
		WT	CT	Metrics		M3SA-Specifics				
				PM	SM	IM	MW	MF		
§IV-B	Reproduce peer-reviewed experiment	WT1	✗	CPU hours	Wh	4	10	median	✗	S1
§IV-C	Analyze M3SA on fundamentally different traces	WT2, WT3	CT1	CPU hours	gCO2	8	1	mean	✓	S2
§IV-D	Adopt M3SA in CO2-aware, dynamic, workload migration	WT2	CT2	CPU hours	tCO2	16	1	median	✗	S3

TABLE II
SYSTEMS UNDER OBSERVATION (SUO).

ID	Source	Scale [#hosts]	Resources per host
S1	SURF	277	128 GB RAM, 16 cores, 2.1 Ghz
S2	Marconi	150	196 GB RAM, 48 cores, 2.1 Ghz
S3	Marconi	2,982	196 GB RAM, 48 cores, 2.1 Ghz

TABLE III

WORKLOAD TRACES USED IN EXPERIMENTS. NAME IS SOURCE AND COLLECTION YEAR (E.G., SURF-22 = SOURCE SURF, YEAR 2022). SUO = SYSTEM UNDER OBSERVATION, SR = SAMPLING RATE, CH = CPU HOURS (MILLIONS), S = SCIENTIFIC, B = BUSINESS-CRITICAL.

ID	Name	Type	SUO	Duration	Jobs	CH	SR
WT1	SURF-22	S	S1	7 days	7,850	0.31	30 s
WT2	Marconi-22	S	S2	30 days	8,316	4.74	20 s
WT3	Solvinty-13	B	S3	30 days	50	0.13	30 s

Solvinty-13 is a business-critical workload hosted in a Solvinty datacenter [48], former BitBrains, a mid-sized cloud service provider in the Netherlands. Solvinty-13 has been used in highly cited, peer-reviewed publications on business-critical workloads [32]. The Solvinty-13 traces long-running jobs, with an average of 2,722 CPU-hours per job.

Carbon traces (public): ENTSO-E Transparency Platform [49] consists of 40 European Transmission System Operators (TSOs), providing open source datasets about energy and sustainability metrics since 2015. In this work, we use traces from 29 countries from the ENTSO-E datasets.

Failure model: The failure model causes infrastructure failures at arbitrary times, with downtime periods following Lsdns04 traces [50], which follow an exponential distribution with known parameters [50]. We assume no checkpointing system, so each job is re-run from the beginning after failure.

Power (draw) models: We only model power draw due to CPU usage in this work. For the Multi-Model (§III-C), we equip the simulator with a library of seven commonly used power-draw models [11], [12], [25], [51], each with a simple analytical formula that links power draw (output) to idle and maximum power draw, and to current utilization; the models include square root, linear, square, cubic, MSE, asymptotic, and asymptotic with cubic attenuation due to DVFS ($P(u) = P_{idle} + \frac{(P_{max}-P_{idle})}{2} \times (1 + u^3 - e^{-u^3/\alpha})$,

TABLE IV

CARBON TRACES USED IN EXPERIMENTS. NAME IS SOURCE, LOCATION OF COLLECTION, AND COLLECTION YEAR (E.G., ENTSOE-NL-22 = SOURCE ENTSOE, COLLECTED IN THE NETHERLANDS, YEAR 2022.)

ID	Name	Location	Duration	Sampling
CT1	ENTSOE-NL-22	The Netherlands	1 year	900 s
CT2	ENTSOE-EU-23	29 European regions	1 year	900 s

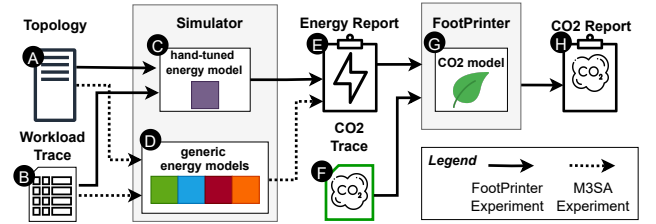


Fig. 8. Overview of experiment 1, adapted and re-run from FootPrinter [16], with M3SA simulation and analysis capabilities.

where u = CPU utilization, P_{idle} and P_{max} are the power used in idle and full capacity states, and α is the utilization fraction at which the host becomes asymptotic). In experiments 1 (§IV-B), 2 (§IV-C), and 3 (§IV-D), we select 4, 8, and 16 combinations of models and specific realistic parameters, respectively. For example, in experiment 2 we use each of the seven power-draw models, and configure only the last with two different configurations, $\alpha = 0.85$ and $\alpha = 1.90$, and $P_{idle} = 32$, $P_{max} = 180$. We detail each power (draw) model in the technical report [38].

B. Peer-reviewed experiment reproduced with M3SA

Niewenhuis et al. propose FootPrinter [16], a tool for predicting the CO2 footprint of datacenters. Figure 8 depicts the experiment we reproduce from [16], leveraging M3SA capabilities to achieve higher accuracy and more detailed explanations; Section IV-A and Table I summarize the design of this experiment. Illustrative for the differences between approaches, whereas FootPrinter employs a single, hand-tuned energy model (C), M3SA uses multiple, generic energy models (D). The other components of the experiment, A, B, and E to H, remain unchanged, enabling a direct comparison between approaches. Our results strongly support MF1.

Figure 9 shows the power draw of the S1 system over one week. Figure 9A shows the results of reproducing the experiment from FootPrinter [16], compared against measured

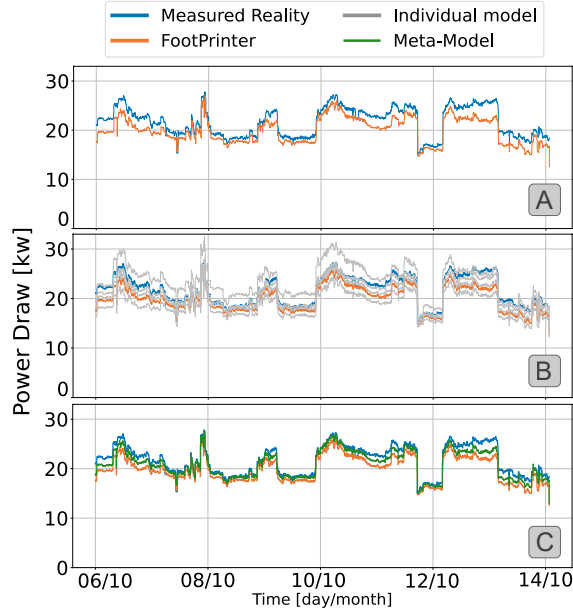


Fig. 9. Simulation results vs. measured reality: (A) FootPrinter simulator, peer-reviewed experiment [16]; (B) Multi-Model with four singular models vs. FootPrinter; (C) Meta-Model derived from the Multi-Model at point B.

reality. Illustrative for a hand-tuned model, the FootPrinter curve approaches in many places the measured reality. As a trade-off, changing the model for a new setting would be complex and error-prone, requiring possibly a similar amount of work to the development of the initial model. (We explore this ability for M3SA, in Section IV-C.)

Better explainability: Figure 9B additionally shows four singular models, whose results were automatically produced through the Multi-Model component. It seems appealing to select the most accurate model, from this sub-plot or peer-reviewed publications, by comparing against the ground truth provided by the measured reality. However, in real-world situations, and increasingly when opaque clouds provide resources, the analyst is often unaware of the ground-truth. The results in Figure 9B show that three of the singular models match the measured reality, while a fourth model constantly overestimates. Using the Multi-Model approach, the analyst can identify singular model biases by contrasting the predictions at each moment in time, thus increasing the overall prediction explainability⁴ over singular-model approaches, where inaccuracies may not even be observable.

“All models are wrong, but some are useful” (George Box, 1976). Simulating with a robustly accurate model is a critical yet non-trivial challenge in ICT. Figure 9C shows the predictions of a Meta-Model, automatically derived from the Multi-Model presented in Figure 9B, and compared against the predictions of the FootPrinter model and the measured reality. We quantify accuracy against measured reality and obtain a MAPE of 7.59% for the average singular model, and a MAPE of 3.81% ($\approx 50\%$ better) for the Meta-Model. The improvement in error rate is due to the *median* aggregation

⁴Resolving the discrepancies between models, and identifying a single best model, requires capabilities beyond the scope of this work.

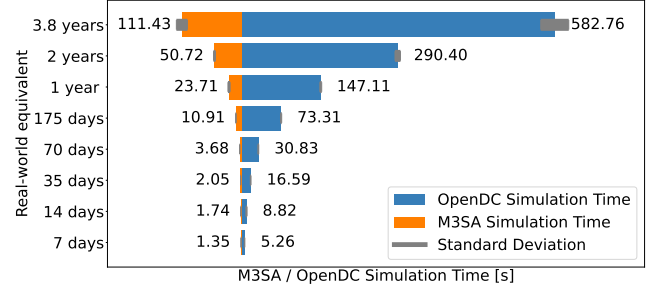


Fig. 10. M3SA performance overhead compared to OpenDC simulation runtime. The plot, which takes its structure from Tukey’s population pyramid, plots on the horizontal axis the system response: (left) M3SA’s overhead, (right) the OpenDC simulation time. On the vertical axis, discretely, is the duration of the workload; higher values cause higher runtimes.

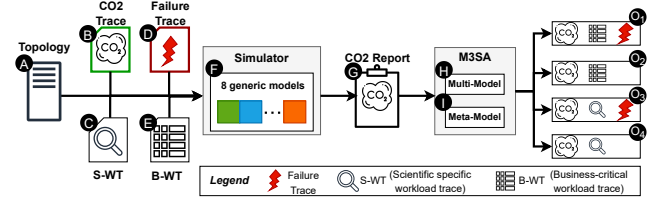


Fig. 11. Experiment 2 overview. Evaluating M3SA on business-critical and scientific-specific workloads, with operational phenomena.

function, which effectively filters extremes, reduces biases, and improves the robustness of the prediction.

Compared to state-of-the-art FootPrinter’s MAPE of 3.15%, the Meta-Model has a higher error rate. However, FootPrinter underlies a manual simulation model, explicitly trained for a single trace (here, SURF-22). In contrast, M3SA enables the versatile, automatic re-use of generic models.

M3SA runtime overhead under realistic conditions: Throughout this experiment, M3SA operated on 20,170 samples, or about 7 days of operation with SURF monitoring rate of 30 seconds [11]. To assess the runtime overhead of M3SA, we further create datasets of various sizes, equivalent to different datacenter operational times, including 2 years (**NFR3**) and up to nearly 4 years, by modifying the monitoring granularity of the simulator.

Figure 10 summarizes the results of this performance-analysis experiment. The M3SA overhead includes synthesizing a Multi-Model, computing the Meta-Model, and outputting the results. We run each experiment 10 times on a standard MacBook M2 Pro and represent the standard deviation in gray; full analysis is available in Appendix B [38]. On average, M3SA adds less than 20% time overhead to the overall simulation process, even for massive-scale datasets, and can simulate years of datacenter operation within minutes (**NFR1**).

C. Evaluating M3SA on fundamentally different traces

Support for NFR2 also introduces the question of which models can be considered. The ability to cope with the diversity of possible workloads is already understood as essential for datacenter simulators [52]. Similarly, failures are known to change how specific workloads are run [52], [53]. We showcase in this section, through an experiment, how M3SA

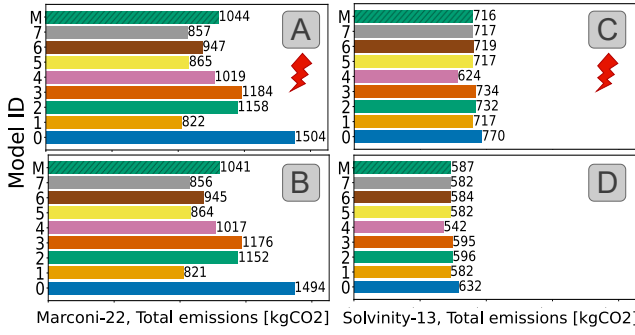


Fig. 12. Total emissions generated by a workload trace running in S2. Left column, Marconi-22 workload: (A) S2 experiences Ldsdn04-based failures, (B) S2 without failures. Right column, Solvinity-13 workload: (C) S2 experiences Ldsdn04-based failures, (D) S2 without failures. Each plot shows the total CO2 emissions predicted by the M3SA Multi-Model using eight generic energy-models (labeled 0-7) and by the Meta-Model (M).

cope with fundamentally different traces and with machine-failures. Section IV-A and Table I summarize the design of this experiment. Figure 11 depicts the experiment design including two kinds of workloads to be run independently (C and E), and the option to include failures (D); these lead to four output combinations, O_1 to O_4 .

Overall, besides the ability of M3SA to provide answers for these complex conditions (MF3), the results depicted in Figure 12 indicate that the presence of failures can add from near-zero (0.28%) to significant (21.9%) CO2 emissions compared to running the workload without system failures, so sustainability studies such as [3] should explicitly include such analysis. Consequently, C-level stakeholders should not make decisions absent such analysis.

We run two fundamentally different traces. Not only do they come from different domains, scientific vs. business-critical, but also, whereas Marconi includes many job arrivals that fit diurnal and day-of-week patterns [31], Solvinity-13 has much longer jobs (see Section IV-A) and forms a stable workload that is largely time-insensitive [32].

Figure 12 shows the environmental impact of the two different workloads, when failures are not or also considered. The results support that M3SA can help identify (overly) biased models, even when considering only the aggregate results. In this experiment, we observe that model 0 (sqrt, introduced in §IV-A, $P(u) = P_{idle} + (P_{max} - P_{idle}) \times \sqrt{u}$) constantly overestimates, predicting up to 54% more than the average of the other models (value computed from Figure 12, subplot A, where model 0 predicts 1,504 kgCO₂, $\approx 54\%$ more than the average of 979 kgCO₂ of models 1-7, value computed as $(1,504 - 979) / 979 \approx 54\%$). This major overestimation, albeit clear in a Multi-Model simulation, would not be observable in a single-model simulation; here, the Meta-Model approach helps.

The reporting features of M3SA allow for tailored evaluation and analysis. For example, when running Marconi-22, the presence of failures leads to only a 0.28% increase in CO2 emissions (Figure 12, sub-plots A and B, value computed as the difference between the meta-model results,

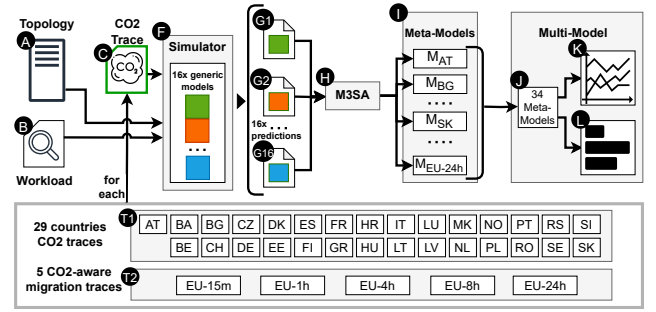


Fig. 13. Overview of experiment 3, assessing CO2 footprint of a workload run in 29 countries, and run using 5 CO2-aware migrations intervals.

labeled M in each sub-plot, normalized by the smallest of the meta-model results, i.e., $(1,044 - 1,041)/1,041$). In contrast, when running Solvinity-13, which has a much longer job duration, failures add 21.9% CO2 emissions, equivalent to 129 kg of CO₂ (Figure 12C and D). Scaling linearly to the scale of Marconi HPC, consisting of 2,982 nodes, operational phenomena would add 2.56 tons of CO₂. M3SA can help C-levels predict such emissions, thus reducing the risk of unexpected emissions of the operated ICT infrastructure and helping operators cope with increasingly tight climate-related regulation.

D. Evaluating M3SA on CO2-aware workload migration

ICT infrastructure accounts for 3% of global greenhouse gases, on par with the aviation industry [54]. To address this issue, C-levels adopt CO2 reporting and instruct their organizations to adopt CO2-aware scheduling and migration policies while still respecting Service Level Agreements [55]; in the upcoming years, regulation could require even more detailed analysis and related actions.

In this section, we analyze how M3SA copes with CO2-aware workload scheduling. Figure 13 depicts the experiment including CO2 traces from 29 countries (T1) and 5 migration traces at different granularities, where the location with the lowest CO2 impact is chosen at given intervals (T2). As summarized by Table I, we equip M3SA with the same 16 singular energy models and ICT infrastructure in each experiment. To enhance prediction robustness and alleviate per-model biases, we synthesize a Meta-Model for each location. Then, we run a greedy CO2-aware migration algorithm and obtain 5 CO2-aware migration traces (T2).

Overall, our results show that CO2-aware workload migration and scheduling processes can help reduce CO2 emissions significantly (MF4) and should be considered more extensively by operators during both capacity planning and operational phases.

We design and prototype a greedy CO2-aware migration algorithm (Appendix C [38], which migrates the workload to the greedy-best CO2 location at time intervals, assuming no migration costs, instant migration, and sufficient infrastructure capacity. For this experiment, we select migration intervals of 15 minutes, 1 hour, 4 hours, 8 hours, and 24 hours (T2).

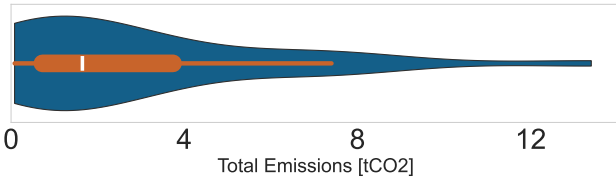


Fig. 14. Total CO2 emitted when running the Marconi-22 workload on S3, assessed for 29 European countries, as of June 2023. The orange box and extended whiskers represent the typical box-and-whiskers plot, overimposed on a violin plot.

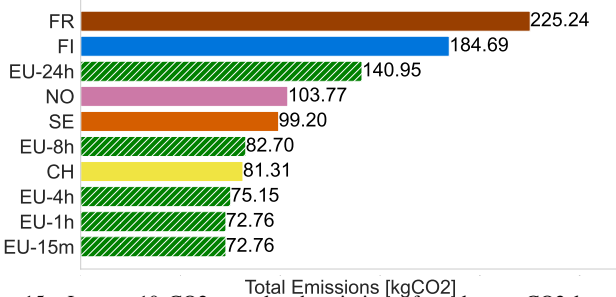


Fig. 15. Lowest 10 CO2 cumulated emissions, from lowest CO2 locations and migration models, at different granularities.

We analyze the number of migrations for each month of 2023 and observe that June has the most overall migrations, while January has the least. We run Marconi-22 workload, spanning over 30 days, in June 2023 and observe the CO2 footprint distribution ranging from less than 100 kg CO2 (e.g., Sweden, Switzerland) up to 13 tons CO2 (e.g., Germany), with an average of 2.92 tons CO2. The differences, some in orders of magnitude, reflect the CO2 intensity of the energy sources; for example, in Switzerland, less than 1% of electricity was obtained from brown sources [56], compared to 37% in Germany [57].

We identify migration intervals of 15 minutes and 1 hour as emitting 11% less CO2 than running the workload integrally in the lowest CO2 location, and 97.5% fewer emissions, equivalent to 2.9 tons than running the workload integrally in an average location; Figure 15 shows top-10 CO2 locations (vertical axis) and cumulated CO2 emissions (horizontal axis). Still, while 24-hour migrations are more sustainable than 26 out of 29 locations, daily migrations can result in 73% more emissions than running the workload in Switzerland, the lowest CO2 location.

The reporting features of M3SA allow for tailored CO2-aware migration and analysis and aid the migration algorithm in selecting the locations by CO2 footprint. This experiment unveils a tradeoff between migration granularity and environmental impact, unveiling new research areas in which M3SA can help identify, compare, and evaluate CO2-aware migrators and schedulers. Such tools would assist datacenter operators in making better-informed decisions about workload scheduling and migration, significantly improving operational footprint.

V. RELATED WORK

Simulating using multiple models is a novelty in computer systems; however, as introduced in Section I, other scientific fields already employ multi-model simulations, such as virology [19], weather and climate simulation [28], or ecology [29].

In ICT, existing simulators currently rely on singular model predictions [11]–[15]. We are the first to propose Multi- and Meta- Model simulation of ICT infrastructure.

Closest to our work, we identify COVID-19 Forecast Hub [19], used to predict the COVID-19 hospitalizations in the United States, and providing 93 models, out of which users can select and leverage in a unified visualization. This is equivalent to the *Multi-Model* component of M3SA. In weather and climate simulation [20], Myhre et al. investigated the geographical distribution of harmful emissions using seven models unified under a singular representation and aggregated (averaged) into a “Mean Model”, helping to contrast biases and improve prediction explainability. This is equivalent to the *Meta-Model* component of M3SA.

Datacenter simulators are useful tools for the community and are widely used in predicting ICT infrastructure. However, no simulator provides multi- and meta-simulation capabilities but relies on singular-model predictions [11]–[15]. Our work adds this capability as an external add-on.

VI. CONCLUSION AND FUTURE WORK

Understanding the performance and climate impact of existing and, more importantly, upcoming datacenters is essential to our society and economy. Although datacenter operators already use simulators to predict infrastructure capabilities, current simulators often rely on singular models, which are insufficient for reliable predictions under diverse scenarios.

In this work, we have designed, implemented, and evaluated M3SA, a framework for multi- and meta-model simulation and analysis of datacenters. We evaluated M3SA using real-world workload and CO2 emission traces, demonstrating M3SA’s ability to predict in complex scenarios, with enhanced explainability and robustness over approaches that only use singular models (even hand-tuned). Such capabilities can aid analysts and C-level decision-makers to better reason about real-world scenarios and make informed decisions. Results show that M3SA’s can simulate years of datacenter operation in minutes, with a 50% lower error rate than singular models, while enabling detailed explanation and without significant overheads. We also showcase how M3SA can simulate system failures and contribute to CO2-aware workload migration and scheduling, reducing CO2 emissions significantly when considering datacenters’ geographical location.

We have released M3SA as an open-sourced system that can be applied as a top layer to datacenter simulators and tested its operation with the commonly used open-source simulator OpenDC. As explained in Section III-E, we envision future research in exploring MCDA, and various methods of leveraging a Meta-Model, such as dynamically allocating weights to predictions of singular models. We also envision future research in employing AI/ML techniques in the Meta-Model aggregation function and are currently embedding M3SA into a simulation-based *digital twin* [58], as part of a major infrastructure project with over 75 partner institutions.

REFERENCES

- [1] G. Andreadis, L. Versluis, F. Mastenbroek, and A. Iosup, "A reference architecture for datacenter scheduling: design, validation, and experiments," in *Proceedings of the International Conference for High Performance Computing, Networking, Storage, and Analysis, SC 2018, Dallas, TX, USA, November 11-16, 2018*, pp. 37:1–37:15, IEEE / ACM, 2018.
- [2] IDC and Equinix, "Idc marketscape – worldwide datacenter services 2023 vendor assessment." <https://www.equinix.nl/resources/analyst-reports/interconnection-colocation-equinix-idc-marketscape>, 2023.
- [3] IDC, "AI datacenter capacity, energy consumption, and carbon emission projections." <https://www.idc.com/getdoc.jsp?containerId=US52131624>, 2024.
- [4] "Predicts 2024: Unraveling tomorrow's cloud computing landscape." <https://www.gartner.com/en/documents/5178763>, 2024.
- [5] "Eu horizon project, graph-massivizer." <https://graph-massivizer.eu/>, 2025.
- [6] C. Ren, D. Wang, B. Urgaonkar, and A. Sivasubramaniam, "Carbon-aware energy capacity planning for datacenters," in *20th IEEE International Symposium on Modeling, Analysis and Simulation of Computer and Telecommunication Systems, MASCOTS 2012, Washington, DC, USA, August 7-9, 2012*, pp. 391–400, IEEE Computer Society, 2012.
- [7] S. Tuli, S. R. Poojara, S. N. Srirama, G. Casale, and N. R. Jennings, "COSCO: container orchestration using co-simulation and gradient based optimization for fog computing environments," *IEEE Trans. Parallel Distributed Syst.*, vol. 33, no. 1, pp. 101–116, 2022.
- [8] R. Prodan, D. Kimovski, A. Bartolini, M. Cochez, A. Iosup, E. Kharlamov, J. Rožanec, L. Vasiliu, and A. L. Vărbănescu, "Towards extreme and sustainable graph processing for urgent societal challenges in europe," in *2022 IEEE Cloud Summit*, pp. 23–30, IEEE, 2022.
- [9] A. Iosup, A. Uta, L. Versluis, G. Andreadis, E. V. Eyk, T. Hegeman, S. Talluri, V. van Beek, and L. Toader, "Massivizing computer systems: A vision to understand, design, and engineer computer ecosystems through and beyond modern distributed systems," in *38th IEEE International Conference on Distributed Computing Systems, ICDCS 2018, Vienna, Austria, July 2-6, 2018*, pp. 1224–1237, IEEE Computer Society, 2018.
- [10] N. Bronson, A. Aghayev, A. Charapko, and T. Zhu, "Metastable failures in distributed systems," in *HotOS '21: Workshop on Hot Topics in Operating Systems, Ann Arbor, Michigan, USA, June, 1-3, 2021* (S. Angel, B. Kasikci, and E. Kohler, eds.), pp. 221–227, ACM, 2021.
- [11] F. Mastenbroek, G. Andreadis, S. Jounaid, W. Lai, J. Burley, J. Bosch, E. V. Eyk, L. Versluis, V. van Beek, and A. Iosup, "Opencd 2.0: Convenient modeling and simulation of emerging technologies in cloud datacenters," in *21st IEEE/ACM International Symposium on Cluster, Cloud and Internet Computing, CCGrid 2021, Melbourne, Australia, May 10-13, 2021* (L. Lefèvre, S. Patterson, Y. C. Lee, H. Shen, S. Ilager, M. Goudarzi, A. N. Toosi, and R. Buyya, eds.), pp. 455–464, IEEE, 2021.
- [12] R. N. Calheiros, R. Ranjan, A. Beloglazov, C. A. F. D. Rose, and R. Buyya, "Cloudsim: a toolkit for modeling and simulation of cloud computing environments and evaluation of resource provisioning algorithms," *Softw. Pract. Exp.*, vol. 41, no. 1, pp. 23–50, 2011.
- [13] T. B. Hewage, S. Ilager, M. A. Rodriguez, and R. Buyya, "Cloudsim express: A novel framework for rapid low code simulation of cloud computing environments," *Softw. Pract. Exp.*, vol. 54, no. 3, pp. 483–500, 2024.
- [14] H. Casanova, "Simgrid: A toolkit for the simulation of application scheduling," in *First IEEE International Symposium on Cluster Computing and the Grid (CCGrid 2001), May 15-18, 2001, Brisbane, Australia*, pp. 430–441, IEEE Computer Society, 2001.
- [15] J. McDonald, J. Dobbs, Y. C. Wong, R. F. da Silva, and H. Casanova, "An exploration of online-simulation-driven portfolio scheduling in workflow management systems," *Future Gener. Comput. Syst.*, vol. 161, pp. 345–360, 2024.
- [16] D. Niewenhuis, S. Talluri, A. Iosup, and T. D. Matteis, "Footprinter: Quantifying data center carbon footprint," in *Companion of the 15th ACM/SPEC International Conference on Performance Engineering, ICPE 2024, London, United Kingdom, May 7-11, 2024* (S. Balsamo, W. J. Knottenbelt, C. L. Abad, and W. Shang, eds.), pp. 189–195, ACM, 2024.
- [17] M. S. Ardebili, A. Acquaviva, L. Benini, and A. Bartolini, "Hazardnet: A thermal hazard prediction framework for datacenters," *Future Gener. Comput. Syst.*, vol. 155, pp. 340–353, 2024.
- [18] B. Aksar, E. Sencan, B. Schwaller, O. Aaziz, V. J. Leung, J. M. Brandt, B. Kulis, M. Egele, and A. K. Coskun, "Runtime performance anomaly diagnosis in production HPC systems using active learning," *IEEE Trans. Parallel Distributed Syst.*, vol. 35, no. 4, pp. 693–706, 2024.
- [19] E. Y. o. Cramer, "The united states covid-19 forecast hub dataset," *Nature*, vol. 9, no. 1, p. 462, 2022.
- [20] G. Myhre *et al.*, "Multi-model simulations of aerosol and ozone radiative forcing due to anthropogenic emission changes during the period 1990–2015," *Atmospheric Chemistry and Physics*, vol. 17, no. 4, pp. 2709–2720, 2017.
- [21] NeilMcAllister, "Google goes dark for 2 minutes, kills 40% of world's net traffic." https://www.theregister.com/2013/08/17/google_outage/.
- [22] S. Moss, "Surgeries and procedures paused at wichita hospitals due to data center outage." Data Center Dynamics, <https://www.datacenterdynamics.com/en/news/surgeries-and-procedures-paused-at-wichita-hospitals-due-to-data-center-outage/>, 2023.
- [23] "Overheating datacenter stopped 2.5 million bank transactions." Laura Dobberstein, The Register, https://www.theregister.com/2023/11/07/overheating_datacenter_singapore/, 2023.
- [24] M. C. S. Filho, R. L. Oliveira, C. C. Monteiro, P. R. M. Inácio, and M. M. Freire, "Cloudsim plus: A cloud computing simulation framework pursuing software engineering principles for improved modularity, extensibility and correctness," in *2017 IFIP/IEEE Symposium on Integrated Network and Service Management (IM), Lisbon, Portugal, May 8-12, 2017*, pp. 400–406, IEEE, 2017.
- [25] R. F. da Silva, A. Orgerie, H. Casanova, R. Tanaka, E. Deelman, and F. Suter, "Accurately simulating energy consumption of i/o-intensive scientific workflows," in *Computational Science - ICCS 2019 - 19th International Conference, Faro, Portugal, June 12-14, 2019, Proceedings, Part I* (J. M. F. Rodrigues, P. J. S. Cardoso, J. M. Monteiro, R. Lam, V. V. Krzhizhanovskaya, M. H. Lees, J. J. Dongarra, and P. M. A. Sloot, eds.), vol. 11536 of *Lecture Notes in Computer Science*, pp. 138–152, Springer, 2019.
- [26] D. Huang, L. Costero, A. Pahlevan, M. Zapater, and D. Atienza, "Cloudprophet: A machine learning-based

- USA, December 7-10, 2003 (S. E. Chick, P. J. Sanchez, D. M. Ferrin, and D. J. Morrice, eds.), pp. 1296–1304, IEEE Computer Society, 2003.
- [37] F. Zarai and P. Nicopolitidis, eds., *Modeling and Simulation of Computer Networks and Systems*. Elsevier, 2015.
- [38] “M3sa technical report, double-blinded.” <https://anonymous.4open.science/r/m3sa/M3SA-technical-report.pdf>.
- [39] The Apache Software Foundation, “Apache parquet,” 2024.
- [40] I. B. Huang, J. Keisler, and I. Linkov, “Multi-criteria decision analysis in environmental sciences: Ten years of applications and trends,” *Science of the total environment*, vol. 409, no. 19, pp. 3578–3594, 2011.
- [41] J.-J. Wang, Y.-Y. Jing, C.-F. Zhang, and J.-H. Zhao, “Review on multi-criteria decision analysis aid in sustainable energy decision-making,” *Renewable and sustainable energy reviews*, vol. 13, no. 9, pp. 2263–2278, 2009.
- [42] A. Ishizaka and P. Nemery, *Multi-criteria decision analysis: methods and software*. John Wiley & Sons, 2013.
- [43] Oracle, “Mape (mean absolute percentage error) documentation.” Oracle Cloud Infrastructure Documentation, https://docs.oracle.com/en/cloud/saas/planning-budgeting-cloud/pfusu/insights_metrics_MAPE.html, 2024. Accessed: 30 April 2025.
- [44] J. J. M. Moreno *et al.*, “Using the r-mape index as a resistant measure of forecast accuracy,” *Psicothema*, 2013.
- [45] T. Chai and R. R. Draxler, “Root mean square error (rmse) or mean absolute error (mae)?—arguments against avoiding rmse in the literature,” *Geoscientific model development*, 2014.
- [46] A. Gunawardana and G. Shani, “A survey of accuracy evaluation metrics of recommendation tasks,” 2009.
- [47] J. L. Herlocker, J. A. Konstan, L. G. Terveen, and J. T. Riedl, “Evaluating collaborative filtering recommender systems,” *ACM Transactions on Information Systems (TOIS)*, vol. 22, no. 1, pp. 5–53, 2004.
- [48] “Solvinity: Secure managed cloud.” Solvinity, 2025.
- [49] European Network of Transmission System Operators for Electricity (ENTSO-E), “Entso-e transparency platform.” Official Website, <https://transparency.entsoe.eu/>, 2025.
- [50] D. Kondo, B. Javadi, A. Iosup, and D. H. J. Epema, “The failure trace archive: Enabling comparative analysis of failures in diverse distributed systems,” in *10th IEEE/ACM International Conference on Cluster, Cloud and Grid Computing, CCGrid 2010, 17-20 May 2010, Melbourne, Victoria, Australia*, pp. 398–407, IEEE Computer Society, 2010.
- [51] X. Fan, W. Weber, and L. A. Barroso, “Power provisioning for a warehouse-sized computer,” in *34th International Symposium on Computer Architecture (ISCA 2007), June 9-13, 2007, San Diego, California, USA* (D. M. Tullsen and B. Calder, eds.), pp. 13–23, ACM, 2007.
- [52] G. Amvrosiadis, J. W. Park, G. R. Ganger, G. A. Gibson, E. Baseman, and N. DeBardeleben, “On the diversity of cluster workloads and its impact on research results,” in *Proceedings of the 2018 USENIX Annual Technical Conference, USENIX ATC 2018, Boston, MA, USA, July 11-13, 2018* (H. S. Gunawi and B. C. Reed, eds.), pp. 533–546, USENIX Association, 2018.
- [53] P. Gill, N. Jain, and N. Nagappan, “Understanding network failures in data centers: measurement, analysis, and implications,” in *Proceedings of the ACM SIGCOMM 2011 Conference on Applications, Technologies, Architectures, and Protocols for Computer Communications, Toronto, ON, Canada, August 15-19, 2011* (S. Keshav, J. Liebeherr, J. W. Byers, and J. C. Mogul, eds.), pp. 350–361, ACM, 2011.
- [54] L. Eeckhout, “FOCAL: A first-order carbon model to assess processor sustainability,” in *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2, ASPLOS 2024, La Jolla, CA, USA, 27 April 2024- 1 May 2024* (R. Gupta, N. B. Abu-Ghazaleh, M. Musuvathi, and D. Tsafir, eds.), pp. 401–415, ACM, 2024.
- [55] A. Souza, S. Jasoria, B. Chakrabarty, A. Bridgwater, A. Lundberg, F. Skogh, A. Ali-Eldin, D. E. Irwin, and P. J. Shenoy, “CASPER: carbon-aware scheduling and provisioning for distributed web services,” in *Proceedings of the 14th International Green and Sustainable Computing Conference, IGSC 2023, Toronto, ON, Canada, October 28-29, 2023*, pp. 67–73, ACM, 2023.
- [56] “Switzerland: Electricity generation, consumption, and co2 emissions.” International Energy Agency (IEA), <https://www.iea.org/countries/switzerland/electricity>, 2025.
- [57] Federal Statistical Office of Germany (Destatis), “Gross electricity production in germany by energy source (2019-2024).” Official Statistics Table, <https://www.destatis.de/EN/Themes/Economic-Sectors-Enterprises/Energy/Production/Tables/gross-electricity-production.html>, 2023.
- [58] N. A. of Engineering, E. National Academies of Sciences, and Medicine, *Foundational Research Gaps and Future Directions for Digital Twins*. Washington, DC: The National Academies Press, 2024.

